

# Deep Autoencoding Gaussian Mixture Model for Unsupervised Anomaly Detection

```
import numpy as np
import pandas as pd
import torch
from data_loader import *
from main import *
from tqdm import tqdm
```

## KDD Cup 1999 Data (10% subset)

This is the data set used for The Third International Knowledge Discovery and Data Mining Tools Competition, which was held in conjunction with KDD-99 The Fifth International Conference on Knowledge Discovery and Data Mining. The competition task was to build a network intrusion detector, a predictive model capable of distinguishing between "bad" connections, called intrusions or attacks, and "good" normal connections. This database contains a standard set of data to be audited, which includes a wide variety of intrusions simulated in a military network environment.

```
#Windows
#transaction_data = pd.read_csv("C:/Users/cncluser/Downloads/ieee-fraud-detection/t
#macOS
transaction_data = pd.read_csv("/Users/nami/Downloads/ieee-fraud-detection/train_tr
transaction_data
```

	TransactionID	isFraud	TransactionDT	TransactionAmt	Product
0	2987000	0	86400	68.50	W
1	2987001	0	86401	29.00	W
2	2987002	0	86469	59.00	W
3	2987003	0	86499	50.00	W

4	2987004	0	86506	50.00	H
...	...	...	...	...	...
590535	3577535	0	15811047	49.00	W
590536	3577536	0	15811049	39.50	W
590537	3577537	0	15811079	30.95	W
590538	3577538	0	15811088	117.00	W
590539	3577539	0	15811131	279.95	W

590540 rows × 394 columns

```
#Windows
#identity_data = pd.read_csv("C:/Users/cncluser/Downloads/ieee-fraud-detection/train_id
#macOS
identity_data = pd.read_csv("/Users/nami/Downloads/ieee-fraud-detection/train_iden
identity_data
```

	TransactionID	id_01	id_02	id_03	id_04	id_05	
0	2987004	0.0	70787.0	NaN	NaN	NaN	
1	2987008	-5.0	98945.0	NaN	NaN	0.0	
2	2987010	-5.0	191631.0	0.0	0.0	0.0	
3	2987011	-5.0	221832.0	NaN	NaN	0.0	
4	2987016	0.0	7460.0	0.0	0.0	1.0	

...	...	...	...	...	...	...	
144228	3577521	-15.0	145955.0	0.0	0.0	0.0	
144229	3577526	-5.0	172059.0	NaN	NaN	1.0	
144230	3577529	-20.0	632381.0	NaN	NaN	-1.0	
144231	3577531	-5.0	55528.0	0.0	0.0	0.0	
144232	3577534	-45.0	339406.0	NaN	NaN	-10.0	

144233 rows × 41 columns

```
data = transaction_data.set_index('TransactionID').join(identity_data.set_index('Tr
data
```

	isFraud	TransactionDT	TransactionAmt	ProductCD	
TransactionID					
2987000	0	86400	68.50	W	1
2987001	0	86401	29.00	W	2
2987002	0	86469	59.00	W	4
2987003	0	86499	50.00	W	1
2987004	0	86506	50.00	H	4

...	...	...	...	...	...
3577535	0	15811047	49.00	W	6
3577536	0	15811049	39.50	W	1
3577537	0	15811079	30.95	W	1
3577538	0	15811088	117.00	W	7
3577539	0	15811131	279.95	W	1

590540 rows × 433 columns

## Pre-processing

"isFraud" = 0 -> normal, "isFraud" = 1 -> anomaly.

Next, the categorical variables are converted to a one hot encoding representation. My implementation is a bit different from the original paper in this aspect. Since I am only using the 10% subset to generate the columns, I get 118 features instead of 120 as reported in the paper.

```

one_hot_ProductCD = pd.get_dummies(data["ProductCD"])
one_hot_card4 = pd.get_dummies(data["card4"])
one_hot_card6 = pd.get_dummies(data["card6"])
one_hot_Pemaildomain = pd.get_dummies(data["P_emaildomain"])
one_hot_Remaildomain = pd.get_dummies(data["R_emaildomain"])
one_hot_M1 = pd.get_dummies(data["M1"])
one_hot_M2 = pd.get_dummies(data["M2"])
one_hot_M3 = pd.get_dummies(data["M3"])
one_hot_M4 = pd.get_dummies(data["M4"])
one_hot_M5 = pd.get_dummies(data["M5"])
one_hot_M6 = pd.get_dummies(data["M6"])
one_hot_M7 = pd.get_dummies(data["M7"])
one_hot_M8 = pd.get_dummies(data["M8"])
one_hot_M9 = pd.get_dummies(data["M9"])
one_hot_id12 = pd.get_dummies(data["id_12"])
one_hot_id15 = pd.get_dummies(data["id_15"])
one_hot_id16 = pd.get_dummies(data["id_16"])
one_hot_id23 = pd.get_dummies(data["id_23"])
one_hot_id27 = pd.get_dummies(data["id_27"])
one_hot_id28 = pd.get_dummies(data["id_28"])
one_hot_id29 = pd.get_dummies(data["id_29"])
one_hot_id30 = pd.get_dummies(data["id_30"])
one_hot_id31 = pd.get_dummies(data["id_31"])

```

```

one_hot_id33 = pd.get_dummies(data["id_33"])
one_hot_id34 = pd.get_dummies(data["id_34"])
one_hot_id35 = pd.get_dummies(data["id_35"])
one_hot_id36 = pd.get_dummies(data["id_36"])
one_hot_id37 = pd.get_dummies(data["id_37"])
one_hot_id38 = pd.get_dummies(data["id_38"])
one_hot_DeviceType = pd.get_dummies(data["DeviceType"])
one_hot_DeviceInfo = pd.get_dummies(data["DeviceInfo"])

```

```

data = data.drop("ProductCD",axis=1)
data = data.drop("card4",axis=1)
data = data.drop("card6",axis=1)
data = data.drop("P_emaildomain",axis=1)
data = data.drop("R_emaildomain",axis=1)
data = data.drop("M1",axis=1)
data = data.drop("M2",axis=1)
data = data.drop("M3",axis=1)
data = data.drop("M4",axis=1)
data = data.drop("M5",axis=1)
data = data.drop("M6",axis=1)
data = data.drop("M7",axis=1)
data = data.drop("M8",axis=1)
data = data.drop("M9",axis=1)
data = data.drop("id_12",axis=1)
data = data.drop("id_15",axis=1)
data = data.drop("id_16",axis=1)
data = data.drop("id_23",axis=1)
data = data.drop("id_27",axis=1)
data = data.drop("id_28",axis=1)
data = data.drop("id_29",axis=1)
data = data.drop("id_30",axis=1)
data = data.drop("id_31",axis=1)
data = data.drop("id_33",axis=1)
data = data.drop("id_34",axis=1)
data = data.drop("id_35",axis=1)
data = data.drop("id_36",axis=1)
data = data.drop("id_37",axis=1)
data = data.drop("id_38",axis=1)
data = data.drop("DeviceType",axis=1)
data = data.drop("DeviceInfo",axis=1)

```

```

data_header = data.columns
data_header = data_header.drop("isFraud")
#data_header

```

```

data = pd.concat([one_hot_ProductCD, one_hot_card4, one_hot_card6,
                  one_hot_Pemaildomain, one_hot_Remaildomain, one_hot_M1,

```

```

one_hot_M2, one_hot_M3, one_hot_M4,
one_hot_M5, one_hot_M6, one_hot_M7,
one_hot_M8, one_hot_M9, one_hot_id12,
one_hot_id15, one_hot_id16, one_hot_id23,
one_hot_id27, one_hot_id28, one_hot_id29,
one_hot_id30, one_hot_id31, one_hot_id33,
one_hot_id34, one_hot_id35, one_hot_id36,
one_hot_id37, one_hot_id38, one_hot_DeviceType,
one_hot_DeviceInfo, data,axis=1)

data.head()

```

	C	H	R	S	W	american express	discover	master
TransactionID								
2987000	0	0	0	0	1	0	1	0
2987001	0	0	0	0	1	0	0	1
2987002	0	0	0	0	1	0	0	0
2987003	0	0	0	0	1	0	0	1
2987004	0	1	0	0	0	0	0	1

5 rows × 2834 columns

```
#data.loc[:, "SAMSUNG SM-G892A Build/NRD90M"]
```

```

proportions = data["isFraud"].value_counts()
print(proportions)
print("Anomaly Percentage", proportions[1] / proportions.sum())

```

```

0    569877
1     20663
Name: isFraud, dtype: int64
Anomaly Percentage 0.03499000914417313

```

```
#proportions_alfa = data["isFraud"].value_counts(normalize=True)
#print(proportions_alfa)
```

Normalize all the numeric variables.

```
cols_to_norm = data_header
print(cols_to_norm)

#data.loc[:, cols_to_norm] = (data[cols_to_norm] - data[cols_to_norm].mean()) / dat
min_cols = data.loc[data["isFraud"]==0 , cols_to_norm].min()
max_cols = data.loc[data["isFraud"]==0 , cols_to_norm].max()

data.loc[:, cols_to_norm] = (data[cols_to_norm] - min_cols) / (max_cols - min_cols)
```

```
Index(['TransactionDT', 'TransactionAmt', 'card1', 'card2', 'card3', 'card5',
      'addr1', 'addr2', 'dist1', 'dist2',
      ...,
      'id_17', 'id_18', 'id_19', 'id_20', 'id_21', 'id_22', 'id_24', 'id_25',
      'id_26', 'id_32'],
      dtype='object', length=401)
```

```
print(min_cols)
print(max_cols)
data.head()
```

```
TransactionDT      86400.000
TransactionAmt      0.251
card1              1000.000
card2              100.000
card3              100.000
...
id_22              10.000
id_24              11.000
id_25              100.000
id_26              100.000
id_32               0.000
Length: 401, dtype: float64
TransactionDT      1.581113e+07
TransactionAmt      3.193739e+04
card1              1.839600e+04
card2              6.000000e+02
card3              2.310000e+02
```

```
...
id_22      4.400000e+01
id_24      2.600000e+01
id_25      5.480000e+02
id_26      2.160000e+02
id_32      3.200000e+01
Length: 401, dtype: float64
```

	C	H	R	S	W	american express	discover	master
TransactionID								
2987000	0	0	0	0	1	0	1	0
2987001	0	0	0	0	1	0	0	1
2987002	0	0	0	0	1	0	0	0
2987003	0	0	0	0	1	0	0	1
2987004	0	1	0	0	0	0	0	1

5 rows × 2834 columns

I saved the preprocessed data into a numpy file format and load it using the pytorch data loader.

```
np.savez_compressed("ieee_fraud", ieee=data.as_matrix())
```

```
/usr/local/var/pyenv/versions/anaconda3-2019.10/envs/ana3env/lib/python3.7/site-packages/ipykernel_launcher.py:1:
"""Entry point for launching an IPython kernel.
```

I initially implemented this to be ran in the command line and use argparse to get the hyperparameters. To make it runnable in a jupyter notebook, I had to create a dummy class for the hyperparameters.

```
class hyperparams():
    def __init__(self, config):
```



```

        self.__dict__.update(**config)
defaults = {
    'lr' : 1e-4,
    'num_epochs' : 200,
    'batch_size' : 1024,
    'gmm_k' : 4,
    'lambda_energy' : 0.1,
    'lambda_cov_diag' : 0.005,
    'pretrained_model' : None,
    'mode' : 'train',
    'use_tensorboard' : False,
    'data_path' : 'ieee_fraud.npz',

    'log_path' : './dagmm/ieee_logs',
    'model_save_path' : './dagmm/ieee_models',
    'sample_path' : './dagmm/ieee_samples',
    'test_sample_path' : './dagmm/ieee_test_samples',
    'result_path' : './dagmm/ieee_results',

    'log_step' : 194//4,
    'sample_step' : 194,
    'model_save_step' : 194,
}

```

```

solver = main(hyperparams(defaults))
accuracy, precision, recall, f_score = solver.test()

```

590540  
2833

0%| | 0/1 [00:00<?, ?it/s]

DaGMM

DaGMM(

```

(encoder): Sequential(
  (0): Linear(in_features=118, out_features=60, bias=True)
  (1): Tanh()
  (2): Linear(in_features=60, out_features=30, bias=True)
  (3): Tanh()
  (4): Linear(in_features=30, out_features=10, bias=True)
  (5): Tanh()
  (6): Linear(in_features=10, out_features=1, bias=True)
)
(decoder): Sequential(
  (0): Linear(in_features=1, out_features=10, bias=True)
  (1): Tanh()
  (2): Linear(in_features=10, out_features=30, bias=True)
  (3): Tanh()
  (4): Linear(in_features=30, out_features=60, bias=True)
)

```

```

    (5): Tanh()
    (6): Linear(in_features=60, out_features=118, bias=True)
)
(estimation): Sequential(
  (0): Linear(in_features=3, out_features=10, bias=True)
  (1): Tanh()
  (2): Dropout(p=0.5, inplace=False)
  (3): Linear(in_features=10, out_features=4, bias=True)
  (4): Softmax(dim=1)
)
)
The number of parameters: 18783

```

---

RuntimeError

Traceback (most recent call last)

```

<ipython-input-16-db7e643ba5ba> in <module>
----> 1 solver = main(hyperparams(defaults))
      2 accuracy, precision, recall, f_score = solver.test()

~/notebook/anaconda_env/dagmm/main.py in main(config)
    23
    24     if config.mode == 'train':
----> 25         solver.train()
    26     elif config.mode == 'test':
    27         solver.test()

~/notebook/anaconda_env/dagmm/solver.py in train(self)
    95         input_data = self.to_var(input_data)
    96
----> 97         total_loss, sample_energy, recon_error, cov_diag = self.dagmm
    98         # Logging
    99         loss = {}

~/notebook/anaconda_env/dagmm/solver.py in dagmm_step(self, input_data)
   160     def dagmm_step(self, input_data):
   161         self.dagmm.train()
--> 162         enc, dec, z, gamma = self.dagmm(input_data)
   163
   164         total_loss, sample_energy, recon_error, cov_diag = self.dagmm.loss_

/usr/local/var/pyenv/versions/anaconda3-2019.10/envs/ana3env/lib/python3.7/site-pac
   539         result = self._slow_forward(*input, **kwargs)
   540     else:
--> 541         result = self.forward(*input, **kwargs)
   542         for hook in self._forward_hooks.values():

```

```

543             hook_result = hook(self, input, result)

~/notebook/anaconda_env/dagmm/model.py in forward(self, x)
    68     def forward(self, x):
    69
--> 70         enc = self.encoder(x)
    71
    72         dec = self.decoder(enc)

/usr/local/var/pyenv/versions/anaconda3-2019.10/envs/ana3env/lib/python3.7/site-pac
539         result = self._slow_forward(*input, **kwargs)
540     else:
--> 541         result = self.forward(*input, **kwargs)
542         for hook in self._forward_hooks.values():
543             hook_result = hook(self, input, result)

/usr/local/var/pyenv/versions/anaconda3-2019.10/envs/ana3env/lib/python3.7/site-pac
    90     def forward(self, input):
    91         for module in self._modules.values():
--> 92             input = module(input)
    93         return input
    94

/usr/local/var/pyenv/versions/anaconda3-2019.10/envs/ana3env/lib/python3.7/site-pac
539         result = self._slow_forward(*input, **kwargs)
540     else:
--> 541         result = self.forward(*input, **kwargs)
542         for hook in self._forward_hooks.values():
543             hook_result = hook(self, input, result)

/usr/local/var/pyenv/versions/anaconda3-2019.10/envs/ana3env/lib/python3.7/site-pac
    85
    86     def forward(self, input):
--> 87         return F.linear(input, self.weight, self.bias)
    88
    89     def extra_repr(self):

/usr/local/var/pyenv/versions/anaconda3-2019.10/envs/ana3env/lib/python3.7/site-pac
1368     if input.dim() == 2 and bias is not None:
1369         # fused op is marginally faster
-> 1370         ret = torch.addmm(bias, input, weight.t())
1371     else:
1372         output = input.matmul(weight.t())

```

**RuntimeError:** size mismatch, m1: [3 x 2833], m2: [118 x 60] at /Users/distiller/pro

I copy pasted the testing code here in the notebook so we could play around the results.

Incrementally compute for the GMM parameters across all training data for a better estimate

```
solver.data_loader.dataset.mode="train"
solver.dagmm.eval()
N = 0
mu_sum = 0
cov_sum = 0
gamma_sum = 0

for it, (input_data, labels) in enumerate(solver.data_loader):
    input_data = solver.to_var(input_data)
    enc, dec, z, gamma = solver.dagmm(input_data)
    phi, mu, cov = solver.dagmm.compute_gmm_params(z, gamma)

    batch_gamma_sum = torch.sum(gamma, dim=0)

    gamma_sum += batch_gamma_sum
    mu_sum += mu * batch_gamma_sum.unsqueeze(-1) # keep sums of the numerator only
    cov_sum += cov * batch_gamma_sum.unsqueeze(-1).unsqueeze(-1) # keep sums of the

    N += input_data.size(0)

train_phi = gamma_sum / N
train_mu = mu_sum / gamma_sum.unsqueeze(-1)
train_cov = cov_sum / gamma_sum.unsqueeze(-1).unsqueeze(-1)

print("N:",N)
print("phi :\n",train_phi)
print("mu :\n",train_mu)
print("cov :\n",train_cov)
```

```
train_energy = []
train_labels = []
train_z = []
for it, (input_data, labels) in enumerate(solver.data_loader):
    input_data = solver.to_var(input_data)
    enc, dec, z, gamma = solver.dagmm(input_data)
    sample_energy, cov_diag = solver.dagmm.compute_energy(z, phi=train_phi, mu=train_mu)

    train_energy.append(sample_energy.data.cpu().numpy())
    train_z.append(z.data.cpu().numpy())
    train_labels.append(labels.numpy())
```

```
train_energy = np.concatenate(train_energy,axis=0)
train_z = np.concatenate(train_z,axis=0)
train_labels = np.concatenate(train_labels,axis=0)
```

## Compute the energy of every sample in the test data

```
solver.data_loader.dataset.mode="test"
test_energy = []
test_labels = []
test_z = []
for it, (input_data, labels) in enumerate(solver.data_loader):
    input_data = solver.to_var(input_data)
    enc, dec, z, gamma = solver.dagmm(input_data)
    sample_energy, cov_diag = solver.dagmm.compute_energy(z, size_average=False)
    test_energy.append(sample_energy.data.cpu().numpy())
    test_z.append(z.data.cpu().numpy())
    test_labels.append(labels.numpy())

test_energy = np.concatenate(test_energy,axis=0)
test_z = np.concatenate(test_z,axis=0)
test_labels = np.concatenate(test_labels,axis=0)
```

```
combined_energy = np.concatenate([train_energy, test_energy], axis=0)
combined_z = np.concatenate([train_z, test_z], axis=0)
combined_labels = np.concatenate([train_labels, test_labels], axis=0)
```

Compute for the threshold energy. Following the paper I just get the highest 20% and treat it as an anomaly. That corresponds to setting the threshold at the 80th percentile.

```
thresh = np.percentile(combined_energy, 100 - 20)
print("Threshold :", thresh)
```

```
pred = (test_energy>thresh).astype(int)
gt = test_labels.astype(int)
```

```
from sklearn.metrics import precision_recall_fscore_support as prf, accuracy_score
```

```
accuracy = accuracy_score(gt, pred)  
precision, recall, f_score, support = prf(gt, pred, average='binary')
```

```
print("Accuracy : {:.4f}, Precision : {:.4f}, Recall : {:.4f}, F-score : {:.4f}")
```

## Visualizing the z space

---

It's a little different from the paper's figure but I assume that's because of the small changes in my implementation.

```
from mpl_toolkits.mplot3d import Axes3D  
import matplotlib.pyplot as plt  
%matplotlib notebook  
fig = plt.figure()  
ax = fig.add_subplot(111, projection='3d')  
ax.scatter(test_z[:,1], test_z[:,0], test_z[:,2], c=test_labels.astype(int))  
ax.set_xlabel('Encoded')  
ax.set_ylabel('Euclidean')  
ax.set_zlabel('Cosine')  
plt.show()
```