

```
import numpy as np
import pandas as pd
import torch
from data_loader import *
from main import *
from tqdm import tqdm
```

```
transaction_data = pd.read_csv("ieee-fraud-detection/train_transaction.csv", header=0)
transaction_data
```

590540 rows × 394 columns

```
identity_data = pd.read_csv("ieee-fraud-detection/train_identity.csv", header=0)
identity_data
```

[illegible]

	TransactionID	id_01	id_02	id_03	id_04	id_05	id_06	id_07	id_08
0	2987004	0.0	70787.0	NaN	NaN	NaN	NaN	NaN	NaN
1	2987008	-5.0	98945.0	NaN	NaN	0.0	-5.0	NaN	NaN
2	2987010	-5.0	191631.0	0.0	0.0	0.0	0.0	NaN	NaN
3	2987011	-5.0	221832.0	NaN	NaN	0.0	-6.0	NaN	NaN
4	2987016	0.0	7460.0	0.0	0.0	1.0	0.0	NaN	NaN
...
144228	3577521	-15.0	145955.0	0.0	0.0	0.0	0.0	NaN	NaN
144229	3577526	-5.0	172059.0	NaN	NaN	1.0	-5.0	NaN	NaN
144230	3577529	-20.0	632381.0	NaN	NaN	-1.0	-36.0	NaN	NaN
144231	3577531	-5.0	55528.0	0.0	0.0	0.0	-7.0	NaN	NaN
144232	3577534	-45.0	339406.0	NaN	NaN	-10.0	-100.0	NaN	NaN

144233 row s × 41 columns

```
data = transaction_data.set_index('TransactionID').join(identity_data.set_index('TransactionID'))
data
```

	isFraud	TransactionDT	TransactionAmt	ProductCD	card1	card2	card3
TransactionID							
2987000	0	86400	68.50	W	13926	NaN	150.0
2987001	0	86401	29.00	W	2755	404.0	150.0
2987002	0	86469	59.00	W	4663	490.0	150.0
2987003	0	86499	50.00	W	18132	567.0	150.0

2987004	0	86506	50.00	H	4497	514.0	150.0
...
3577535	0	15811047	49.00	W	6550	NaN	150.0
3577536	0	15811049	39.50	W	10444	225.0	150.0
3577537	0	15811079	30.95	W	12037	595.0	150.0
3577538	0	15811088	117.00	W	7826	481.0	150.0
3577539	0	15811131	279.95	W	15066	170.0	150.0

590540 rows × 433 columns

Pre-processing

"isFraud" = 0 -> normal, "isFraud" = 1 -> anomaly.

Next, the categorical variables are converted to a one hot encoding representation. My implementation is a bit different from the original paper in this aspect. Since I am only using the 10% subset to generate the columns, I get 118 features instead of 120 as reported in the paper.

```
one_hot_ProductCD = pd.get_dummies(data["ProductCD"])
one_hot_card4 = pd.get_dummies(data["card4"])
one_hot_card6 = pd.get_dummies(data["card6"])
one_hot_Pemaildomain = pd.get_dummies(data["P_emaildomain"])
one_hot_Remaildomain = pd.get_dummies(data["R_emaildomain"])
one_hot_M1 = pd.get_dummies(data["M1"])
one_hot_M2 = pd.get_dummies(data["M2"])
one_hot_M3 = pd.get_dummies(data["M3"])
one_hot_M4 = pd.get_dummies(data["M4"])
one_hot_M5 = pd.get_dummies(data["M5"])
one_hot_M6 = pd.get_dummies(data["M6"])
one_hot_M7 = pd.get_dummies(data["M7"])
one_hot_M8 = pd.get_dummies(data["M8"])
one_hot_M9 = pd.get_dummies(data["M9"])
one_hot_id12 = pd.get_dummies(data["id_12"])
one_hot_id15 = pd.get_dummies(data["id_15"])
one_hot_id16 = pd.get_dummies(data["id_16"])
one_hot_id23 = pd.get_dummies(data["id_23"])
one_hot_id27 = pd.get_dummies(data["id_27"])
one_hot_id28 = pd.get_dummies(data["id_28"])
one_hot_id29 = pd.get_dummies(data["id_29"])
one_hot_id30 = pd.get_dummies(data["id_30"])
one_hot_id31 = pd.get_dummies(data["id_31"])
one_hot_id33 = pd.get_dummies(data["id_33"])
one_hot_id34 = pd.get_dummies(data["id_34"])
one_hot_id35 = pd.get_dummies(data["id_35"])
one_hot_id36 = pd.get_dummies(data["id_36"])
one_hot_id37 = pd.get_dummies(data["id_37"])
one_hot_id38 = pd.get_dummies(data["id_38"])
one_hot_DeviceType = pd.get_dummies(data["DeviceType"])
one_hot_DeviceInfo = pd.get_dummies(data["DeviceInfo"])
```

```
data = data.drop("ProductCD",axis=1)
data = data.drop("card4",axis=1)
data = data.drop("card6",axis=1)
data = data.drop("P_emaildomain",axis=1)
```

```
data = data.drop("R_emaildomain",axis=1)
data = data.drop("M1",axis=1)
data = data.drop("M2",axis=1)
data = data.drop("M3",axis=1)
data = data.drop("M4",axis=1)
data = data.drop("M5",axis=1)
data = data.drop("M6",axis=1)
data = data.drop("M7",axis=1)
data = data.drop("M8",axis=1)
data = data.drop("M9",axis=1)
data = data.drop("id_12",axis=1)
data = data.drop("id_15",axis=1)
data = data.drop("id_16",axis=1)
data = data.drop("id_23",axis=1)
data = data.drop("id_27",axis=1)
data = data.drop("id_28",axis=1)
data = data.drop("id_29",axis=1)
data = data.drop("id_30",axis=1)
data = data.drop("id_31",axis=1)
data = data.drop("id_33",axis=1)
data = data.drop("id_34",axis=1)
data = data.drop("id_35",axis=1)
data = data.drop("id_36",axis=1)
data = data.drop("id_37",axis=1)
data = data.drop("id_38",axis=1)
data = data.drop("DeviceType",axis=1)
data = data.drop("DeviceInfo",axis=1)
```

```
data = pd.concat([one_hot_ProductCD, one_hot_card4, one_hot_card6,
                  one_hot_Pemaildomain, one_hot_Remaildomain, one_hot_M1,
                  one_hot_M2, one_hot_M3, one_hot_M4,
                  one_hot_M5, one_hot_M6, one_hot_M7,
                  one_hot_M8, one_hot_M9, one_hot_id12,
                  one_hot_id15, one_hot_id16, one_hot_id23,
                  one_hot_id27, one_hot_id28, one_hot_id29,
                  one_hot_id30, one_hot_id31, one_hot_id33,
                  one_hot_id34, one_hot_id35, one_hot_id36,
                  one_hot_id37, one_hot_id38, one_hot_DeviceType,
                  one_hot_DeviceInfo, data],axis=1)

data.head()
```

	C	H	R	S	W	american express	discover	mastercard	visa	charge card
TransactionID										
2987000	0	0	0	0	1	0	1	0	0	0
2987001	0	0	0	0	1	0	0	1	0	0
2987002	0	0	0	0	1	0	0	0	1	0
2987003	0	0	0	0	1	0	0	1	0	0
2987004	0	1	0	0	0	0	0	1	0	0

5 row s × 2834 columns

```
#data.loc[:, "SAMSUNG SM-G892A Build/NRD90M"]
```

```
TransactionID
2987000    0
2987001    0
2987002    0
2987003    0
2987004    1
..
3577535    0
3577536    0
3577537    0
3577538    0
3577539    0
Name: SAMSUNG SM-G892A Build/NRD90M, Length: 590540, dtype: uint8
```

```
proportions = data["isFraud"].value_counts()
print(proportions)
print("Anomaly Percentage",proportions[1] / proportions.sum())
```

```
0    569877
1     20663
Name: isFraud, dtype: int64
Anomaly Percentage 0.03499000914417313
```

```
#proportions_alfa = data["type"].value_counts(normalize=True)
#print(proportions_alfa)
```

Normalize all the numeric variables.

```
cols_to_norm = ["duration", "src_bytes", "dst_bytes", "wrong_fragment", "urgent",
                "hot", "num_failed_logins", "num_compromised", "num_root",
                "num_file_creations", "num_shells", "num_access_files", "count", "srv_count",
                "serror_rate", "srv_serror_rate", "rerror_rate", "srv_rerror_rate", "same_srv_rate",
                "diff_srv_rate", "srv_diff_host_rate", "dst_host_count", "dst_host_srv_count", "dst_host_same_srv_rate",
                "dst_host_diff_srv_rate", "dst_host_same_src_port_rate", "dst_host_srv_diff_host_rate",
                "dst_host_serror_rate", "dst_host_srv_serror_rate", "dst_host_rerror_rate", "dst_host_srv_rerror_rate" ]

#data.loc[:, cols_to_norm] = (data[cols_to_norm] - data[cols_to_norm].mean()) / data[cols_to_norm].std()
min_cols = data.loc[data["type"]==0, cols_to_norm].min()
max_cols = data.loc[data["type"]==0, cols_to_norm].max()

data.loc[:, cols_to_norm] = (data[cols_to_norm] - min_cols) / (max_cols - min_cols)
```

I saved the preprocessed data into a numpy file format and load it using the pytorch data loader.

```
np.savez_compressed("kdd_cup",kdd=data.as_matrix())
```

```
C:\Users\cncluser\Anaconda3\lib\site-packages\ipykernel_launcher.py:1: FutureWarning: Method .as_matrix will be removed in
    """Entry point for launching an IPython kernel.
```

I initially implemented this to be ran in the command line and use argparse to get the hyperparameters. To make it runnable in a jupyter notebook, I had to create a dummy class for the hyperparameters.

```

class hyperparams():
    def __init__(self, config):
        self.__dict__.update(**config)
defaults = {
    'lr' : 1e-4,
    'num_epochs' : 200,
    'batch_size' : 1024,
    'gmm_k' : 4,
    'lambda_energy' : 0.1,
    'lambda_cov_diag' : 0.005,
    'pretrained_model' : None,
    'mode' : 'train',
    'use_tensorboard' : False,
    'data_path' : 'kdd_cup.npz',

    'log_path' : './dagmm/logs',
    'model_save_path' : './dagmm/models',
    'sample_path' : './dagmm/samples',
    'test_sample_path' : './dagmm/test_samples',
    'result_path' : './dagmm/results',

    'log_step' : 194//4,
    'sample_step' : 194,
    'model_save_step' : 194,
}


```

```

solver = main(hyperparams(defaults))
accuracy, precision, recall, f_score = solver.test()

```

Elapsed 0:00:13.076672/0:00:07.943268 -- 0:34:43.567725 , Epoch [2/200], Iter [48/194], lr 0.0001, total_loss: 0.1430, sam



png

```

phi tensor([0.2356, 0.1980, 0.3282, 0.2383]) mu tensor([[ -0.5328,  0.8382,  0.5090],
[ -0.4987,  0.8608,  0.4772],
[ -0.4941,  0.8639,  0.4729],
[ -0.5665,  0.8157,  0.5406]]) cov tensor([[[ 0.2696,  0.1789, -0.2516],
[ 0.1789,  0.1193, -0.1676],
[ -0.2516, -0.1676,  0.2355]],

[[ 0.2868,  0.1903, -0.2676],
[ 0.1903,  0.1268, -0.1781],
[ -0.2676, -0.1781,  0.2504]],

[[ 0.2885,  0.1915, -0.2693],
[ 0.1915,  0.1275, -0.1792],
[ -0.2693, -0.1792,  0.2519]],

[[ 0.2497,  0.1658, -0.2331],
[ 0.1658,  0.1105, -0.1553],
[ -0.2331, -0.1553,  0.2182]]])

```

29%

| 57/194 [00:03<00:07, 17.44it/s]

KeyboardInterrupt

Traceback (most recent call last)

```

<ipython-input-10-db7e643ba5ba> in <module>
----> 1 solver = main(hyperparams(defaults))

```

```

2 accuracy, precision, recall, f_score = solver.test()

~\dagmm\main.py in main(config)
23
24     if config.mode == 'train':
--> 25         solver.train()
26     elif config.mode == 'test':
27         solver.test()

~\dagmm\solver.py in train(self)
95         input_data = self.to_var(input_data)
96
--> 97         total_loss, sample_energy, recon_error, cov_diag = self.dagmm_step(input_data)
98         # Logging
99         loss = {}

~\dagmm\solver.py in dagmm_step(self, input_data)
162     enc, dec, z, gamma = self.dagmm(input_data)
163
--> 164     total_loss, sample_energy, recon_error, cov_diag = self.dagmm.loss_function(input_data, dec, z, gamma, sel
165
166     self.reset_grad()

~\dagmm\model.py in loss_function(self, x, x_hat, z, gamma, lambda_energy, lambda_cov_diag)
166     phi, mu, cov = self.compute_gmm_params(z, gamma)
167
--> 168     sample_energy, cov_diag = self.compute_energy(z, phi, mu, cov)
169
170     loss = recon_error + lambda_energy * sample_energy + lambda_cov_diag * cov_diag

~\dagmm\model.py in compute_energy(self, z, phi, mu, cov, size_average)
121     k, D, _ = cov.size()
122
--> 123     z_mu = (z.unsqueeze(1) - mu.unsqueeze(0))
124
125     cov_inverse = []

KeyboardInterrupt:

```

I copy pasted the testing code here in the notebook so we could play around the results.

Incrementally compute for the GMM parameters across all training data for a better estimate

```

solver.data_loader.dataset.mode="train"
solver.dagmm.eval()
N = 0
mu_sum = 0
cov_sum = 0
gamma_sum = 0

for it, (input_data, labels) in enumerate(solver.data_loader):
    input_data = solver.to_var(input_data)
    enc, dec, z, gamma = solver.dagmm(input_data)
    phi, mu, cov = solver.dagmm.compute_gmm_params(z, gamma)

    batch_gamma_sum = torch.sum(gamma, dim=0)

    gamma_sum += batch_gamma_sum

```

```

mu_sum += mu * batch_gamma_sum.unsqueeze(-1) # keep sums of the numerator only
cov_sum += cov * batch_gamma_sum.unsqueeze(-1).unsqueeze(-1) # keep sums of the numerator only

N += input_data.size(0)

train_phi = gamma_sum / N
train_mu = mu_sum / gamma_sum.unsqueeze(-1)
train_cov = cov_sum / gamma_sum.unsqueeze(-1).unsqueeze(-1)

print("N:",N)
print("phi :\n",train_phi)
print("mu :\n",train_mu)
print("cov :\n",train_cov)

```

```

train_energy = []
train_labels = []
train_z = []
for it, (input_data, labels) in enumerate(solver.data_loader):
    input_data = solver.to_var(input_data)
    enc, dec, z, gamma = solver.dagmm(input_data)
    sample_energy, cov_diag = solver.dagmm.compute_energy(z, phi=train_phi, mu=train_mu, cov=train_cov, size_average=False)

    train_energy.append(sample_energy.data.cpu().numpy())
    train_z.append(z.data.cpu().numpy())
    train_labels.append(labels.numpy())

train_energy = np.concatenate(train_energy,axis=0)
train_z = np.concatenate(train_z,axis=0)
train_labels = np.concatenate(train_labels,axis=0)

```

Compute the energy of every sample in the test data

```

solver.data_loader.dataset.mode="test"
test_energy = []
test_labels = []
test_z = []
for it, (input_data, labels) in enumerate(solver.data_loader):
    input_data = solver.to_var(input_data)
    enc, dec, z, gamma = solver.dagmm(input_data)
    sample_energy, cov_diag = solver.dagmm.compute_energy(z, size_average=False)
    test_energy.append(sample_energy.data.cpu().numpy())
    test_z.append(z.data.cpu().numpy())
    test_labels.append(labels.numpy())

test_energy = np.concatenate(test_energy,axis=0)
test_z = np.concatenate(test_z,axis=0)
test_labels = np.concatenate(test_labels,axis=0)

```

```

combined_energy = np.concatenate([train_energy, test_energy], axis=0)
combined_z = np.concatenate([train_z, test_z], axis=0)
combined_labels = np.concatenate([train_labels, test_labels], axis=0)

```

Compute for the threshold energy. Following the paper I just get the highest 20% and treat it as an anomaly. That corresponds to setting the threshold at the 80th percentile.


```
thresh = np.percentile(combined_energy, 100 - 20)
print("Threshold :", thresh)
```

```
pred = (test_energy>thresh).astype(int)
gt = test_labels.astype(int)
```

```
from sklearn.metrics import precision_recall_fscore_support as prf, accuracy_score
```

```
accuracy = accuracy_score(gt,pred)
precision, recall, f_score, support = prf(gt, pred, average='binary')
```

```
print("Accuracy : {:.4f}, Precision : {:.4f}, Recall : {:.4f}, F-score : {:.4f}".format(accuracy,precision, recall, f_
```



Visualizing the z space

It's a little different from the paper's figure but I assume that's because of the small changes in my implementation.

```
from mpl_toolkits.mplot3d import Axes3D
import matplotlib.pyplot as plt
%matplotlib notebook
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(test_z[:,1],test_z[:,0], test_z[:,2], c=test_labels.astype(int))
ax.set_xlabel('Encoded')
ax.set_ylabel('Euclidean')
ax.set_zlabel('Cosine')
plt.show()
```